

# STAT 3022 LAB WEEK 2

## Loading Data Frame into R

Example 1: Climate of Twin Cities in Summer 2012

7. read.table – Load data frames into R

Download the file summer2012.csv from this course website to your machine, then load it into R.

```
A <- read.table("C:/Users/new/Desktop/STAT30222020spring/LAB1/summer2012.csv",
               header = TRUE, sep = ",")
#Apple: cmd+opt+c; Windows two backslash or slash
```

9. mean – What is the mean temperatures of summer 2012?

```
mean(A$AVG)
```

```
## [1] 75.04348
```

The answer is 75.04 .

9\*. tapply(vector, index, function)

```
tapply(A$AVG, A$MN, mean)
```

```
##          6          7          8
## 72.50000 80.32258 72.22581
```

10. sd and var – How about the standard deviation and variance of the daily temperature in summer 2012?

```
sd(A$AVG)
```

```
## [1] 6.717065
```

```
var(A$AVG)
```

```
## [1] 45.11897
```

11. max and min – What are the warmest and coldest daily average temperatures, respectively?

```
max(A$AVG)
```

```
## [1] 91
```

```
min(A$AVG)
```

```
## [1] 59
```

To find the locations that the maximum and minimum occur, use which.max and which.min, respectively.

```
which.max(A$AVG)
```

```
## [1] 34
```

the maximum occurred on the 34th day of summer, e.g. July 4th.

```
which.min(A$AVG)
```

```
## [1] 1
```

the minimum occurred on the 1st day of summer, e.g. June 1st.

To calculate pairwise maximums and minimums, use `pmax` and `pmin`, respectively.

```
pmax(c(2, 3, 4), c(5, 2, 6))
```

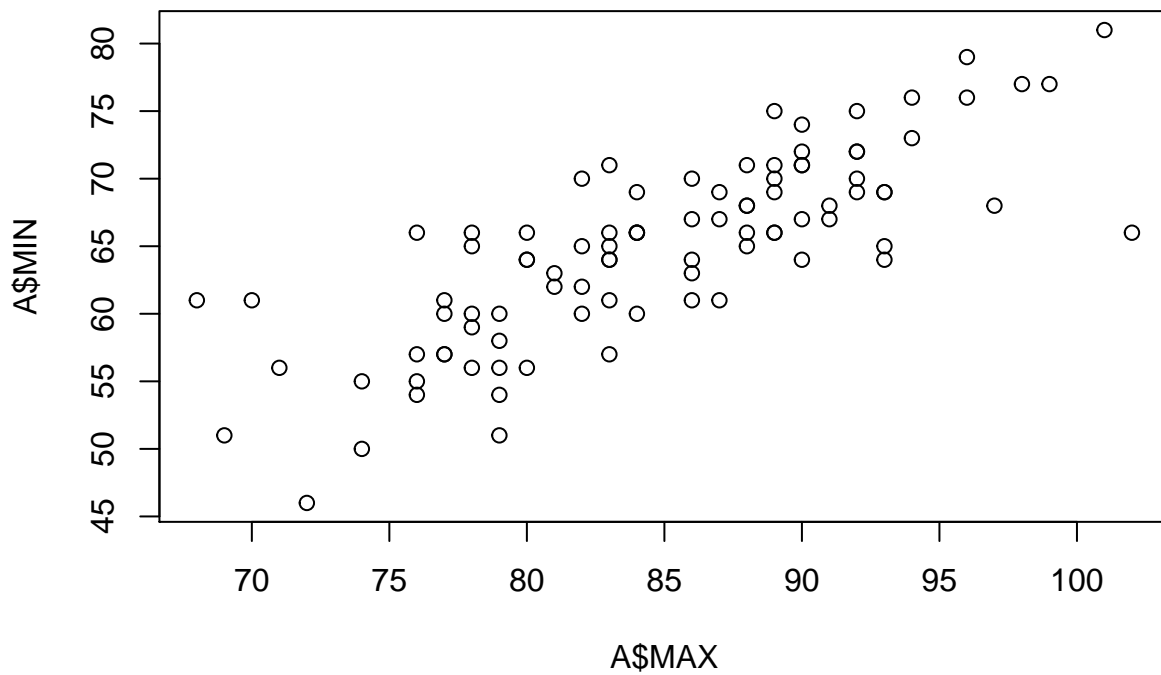
```
## [1] 5 3 6
```

```
pmin(c(2, 3, 4), c(5, 2, 6))
```

```
## [1] 2 2 4
```

12. `plot` – Scatter plot of maximum vs minimum temperatures

```
plot(A$MAX, A$MIN)
```




Scatter plot of the average temperature by day

13. `lines` – Connect the dots on the most recently drawn scatter plot, making it a line chart

```
plot(A$AVG)
```

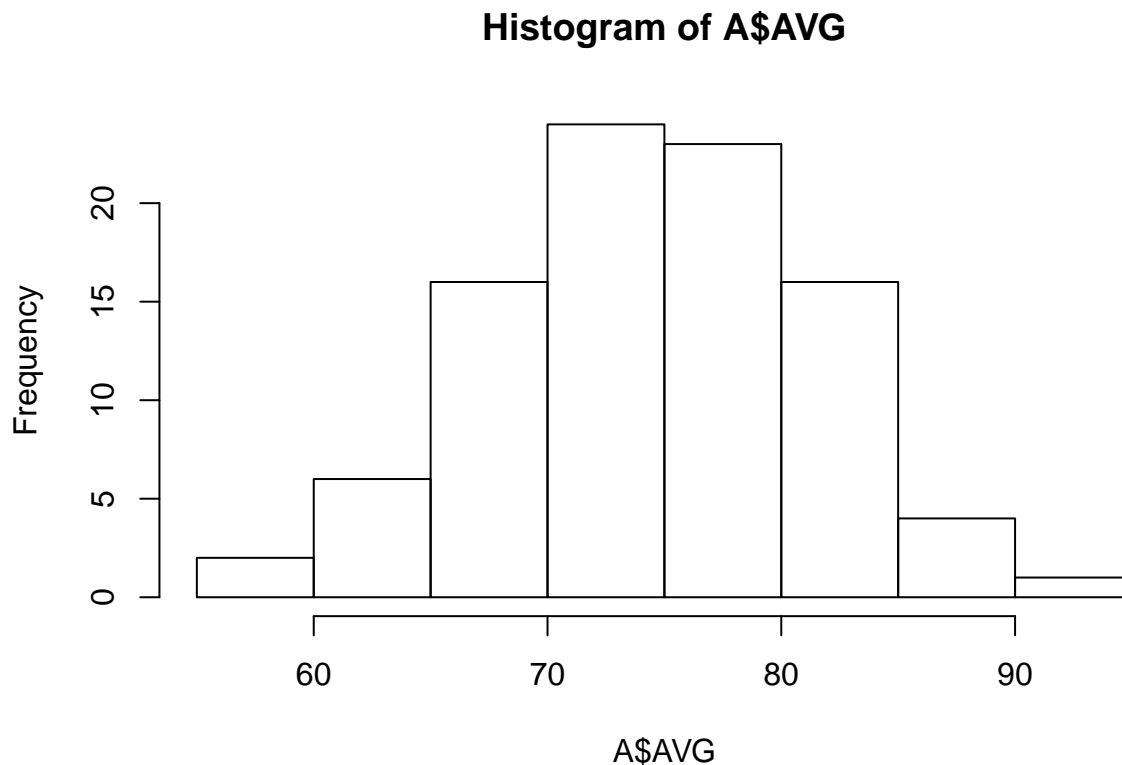
```
lines(A$AVG)
```



Note: Do not turn off the scatter plot before executing the lines command, or it would not work!

14. hist – Plotting histogram for the daily average temperature

```
hist(A$AVG)
```



## SIMPLE LINEAR REGRESSION - AN EXAMPLE USING R

### Example 1: Food Expenditure versus Income


```
food=read.table( "C:/Users/new/Desktop/STAT30222020spring/LAB2/food.csv", sep = ",",header = TRUE)
head(food)
```

```
##   food_exp income
## 1   115.22   3.69
## 2   135.98   4.39
## 3   119.34   4.75
## 4   114.96   6.03
## 5   187.05  12.47
## 6   243.92  12.98
```

It is always a good idea to visually inspect the data in a scatter diagram, which can be created using the function `plot()`. Figure 2.2 is a scatter diagram of food expenditure on income, suggesting that there is a positive relationship between income and food expenditure.

```
plot(food$income, food$food_exp,
     ylim=c(0, max(food$food_exp)),
     xlim=c(0, max(food$income)),
     xlab="weekly income in $100",
     ylab="weekly food expenditure in $",
     main="Figure 2.2: A scatter diagram for the food expenditure model"
)
```

**Figure 2.2: A scatter diagram for the food expenditure model**



### Estimating a Linear Regression

The R function for estimating a linear regression model is `lm(y~x, data)` which, used just by itself does not show any output; It is useful to give the model a name, such as `mod1`, then show the results using `summary(mod1)`. If you are interested in only some of the results of the regression, such as the estimated coefficients, you can retrieve them using specific functions, such as the function `coef()`. For the food expenditure data, the regression model will be

$$\text{food\_exp} = \beta_0 + \beta_1 \text{income} + \epsilon$$

```
mod1 <- lm(food_exp ~ income, data = food)
smod1 <- summary(mod1)
smod1
```

```
##
## Call:
## lm(formula = food_exp ~ income, data = food)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -223.025  -50.816   -6.324   67.879  212.044
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   83.416     43.410   1.922  0.0622 .
## income       10.210      2.093   4.877 1.95e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 89.52 on 38 degrees of freedom
## Multiple R-squared:  0.385, Adjusted R-squared:  0.3688
## F-statistic: 23.79 on 1 and 38 DF,  p-value: 1.946e-05

b1 <- coef(mod1)[[1]]
b2 <- coef(mod1)[[2]]
```

The function `coef()` returns a list containing the estimated coefficients, where a specific coefficient can be accessed by its position in the list. For example, the estimated value of  $\beta_0$  is `b1 <- coef(mod1)[[1]]`, which is equal to 83.416002, and the estimated value of  $\beta_1$  is `b2 <- coef(mod1)[[2]]`, which is equal to 10.209643.

The intercept parameter,  $\beta_0$ , is usually of little importance in econometric models; we are mostly interested in the slope parameter,  $\beta_1$ . The estimated value of  $\beta_1$  suggests that the food expenditure for an average family increases by 10.209643 when the family income increases by 1 unit, which in this case is \$100. The R function `abline()` adds the regression line to the previously plotted scatter diagram, as Figure 2.3 shows.

```
plot(food_exp ~ income, data = food)
abline(mod1)
```




Figure 2.3: Scatter diagram and regression line for the food expenditure model How can one retrieve various regression results? These results exist in two R objects produced by the `lm()` function: the regression object, such as `mod1` in the above code sequence, and the regression summary, which I denoted by `smod1`. The next code shows how to list the names of all results in each object.

```
names(mod1)

## [1] "coefficients" "residuals"    "effects"      "rank"
## [5] "fitted.values" "assign"        "qr"           "df.residual"
```

```
## [9] "xlevels"          "call"          "terms"          "model"
```

```
names(smod1)
```

```
## [1] "call"          "terms"          "residuals"      "coefficients"
## [5] "aliased"        "sigma"          "df"              "r.squared"
## [9] "adj.r.squared" "fstatistic"     "cov.unscaled"
```

To retrieve a particular result you just refer to it with the name of the object, followed by the \$ sign and the name of the result you wish to retrieve. For instance, if we want the vector of coefficients from mod1, we refer to it as mod1\$coefficients and smod1\$coefficients:

```
mod1$coefficients
```

```
## (Intercept)      income
##      83.41600      10.20964
```

```
smod1$coefficients
```

```
##           Estimate Std. Error t value    Pr(>|t|)
## (Intercept) 83.41600  43.410163  1.921578 6.218242e-02
## income      10.20964   2.093264  4.877381 1.945862e-05
```

As we have seen before, however, some of these results can be retrieved using specific functions, such as coef(mod1), resid(mod1), fitted(mod1), and vcov(mod1).

## Example 2


```
x = c(21,34,6,47,10,49,23,32,12,16,29,49,28,8,57,9,31,10,21,
      26,31,52,21,8,18,5,18,26,27,26,32,2,59,58,19,14,16,9,23,
      28,34,70,69,54,39,9,21,54,26)
y = c(47,76,33,78,62,78,33,64,83,67,61,85,46,53,55,71,59,41,82,
      56,39,89,31,43,29,55, 81,82,82,85,59,74,80,88,29,58,71,60,
      86,91,72,89,80,84,54,71,75,84,79)
reg <- lm(y~x) #Create a linear model
summary(reg)
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -32.641  -8.865   3.449  12.438  25.213
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  52.8889     4.4944  11.768 1.3e-15 ***
## x             0.4606     0.1351   3.411 0.00134 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 16.57 on 47 degrees of freedom
## Multiple R-squared:  0.1984, Adjusted R-squared:  0.1814
## F-statistic: 11.63 on 1 and 47 DF,  p-value: 0.00134
```

```
resid(reg) #List of residuals
```


```
##          1          2          3          4          5
## -15.56228406  7.44941046 -22.65270081  3.46110498  4.50474366
##          6          7          8          9         10
##  2.53982722 -30.48356182 -3.62931177  24.58346589  6.74091036
##         11         12         13         14         15
## -5.24739512  9.53982722 -19.78675624 -3.57397858 -24.14528385
##         16         17         18         19         20
## 13.96538254 -8.16867289 -16.49525634 19.43771594 -8.86547847
##         21         22         23         24         25
## -28.16867289 12.15791057 -31.56228406 -13.57397858 -32.18036741
##         26         27         28         29         30
## -0.19206193 19.81963259 17.13452153 16.67388265 20.13452153
##         31         32         33         34         35
## -8.62931177 20.18985472 -0.06656161  8.39407727 -32.64100629
##         36         37         38         39         40
## -1.33781187 10.74091036  2.96538254 22.51643818 25.21324376
##         41         42         43         44         45
##  3.44941046  3.86641067 -4.67295044  6.23663280 -16.85378395
##         46         47         48         49
## 13.96538254 12.43771594  6.23663280 14.13452153
```

```
plot(density(resid(reg))) #A density plot
```




```
qqnorm(resid(reg)) # A quantile normal plot - good for checking normality
qqline(resid(reg))
```





```
boxplot(resid(reg)) #A Boxplot plot
```



```
hist(resid(reg)) #A Histogram plot
```

**Histogram of resid(reg)**

